Kaggle-1: DeepTabular

- **Instructions:**
  1. You may use any library, toolkit, architecture design, loss function, optimizer, or training paradigm of your choice.
  2. You may select any training metric you prefer; however, evaluation will be conducted solely using **Mean Squared Error (MSE)**.
  3. The use of pre-trained models for any sub-task or task is strictly prohibited. Violation of this rule will result in a score of **0 marks** for the entire competition.
  4. Data preprocessing on the provided dataset is not permitted. Any such modifications will result in a score of **0 marks** for the entire competition.
  5. You will be competing against one another. There are n possible methods for developing your solution's **architecture, optimizer, and training paradigm**. If three or more submissions are found to have architectures that are **60% or more similar**, those submissions will receive **0 marks** and may be reported to DAC for plagiarism. (Note: Similarity testing will be automated.)
  6. If the complete pipelines of any two submissions are found to be similar, both will receive **0 marks** and will be reported to DAC for plagiarism.
  7. You must submit your code using the stencil that will be provided two days prior to the submission deadline.
- **Dataset:** The dataset provided is clean and unambiguous. It is a tabular dataset containing:
  - **3 Independent Variables** [F1, F2, F3]
  - **1 Dependent Variable** [OUT]
  - **1 Identifier Variable** [INDEX]
- **Task:**
  - **Task:** Develop a model that accepts three input variables (each of type *int*) and produces a single output variable (of type *float*).
  - Ensure your submission meets the following requirements. (Hint: Save your CSV file with `index=False`.)
    1. The first column must be labeled **ID** and the second column **OUT**.
    2. The **ID** column should contain the corresponding IDs (INDEX) for each row in the test set.
    3. The **OUT** column should contain your prediction values.

Competition Link